


Machine Learning Overview: CPU vs GPU vs Multi-GPUs Performance

Sonal Rami

23/11/2023


Machine learning tools



Software tools and libraries	Hardware
<ul style="list-style-type: none">• TensorFlow• Keras• Scikit-learn• Pandas• NumPy• Jupyter Notebooks• Python	<p>CPUs: Essential for general-purpose tasks</p> <p>GPUs: Provide significant speedup for matrix calculations and ML computations, and deep learning tasks</p> <p>TPUs (Tensor Processing Units)</p>

- Single GPU can significantly speed up machine learning tasks compared to CPU
- Multi-GPU provides enhanced capabilities for handling very large datasets and complex models

TensorFlow for Single GPU



1. Install Prerequisites:

Ensure you have a compatible NVIDIA GPU

Install NVIDIA drivers, CUDA Toolkit, and cuDNN

(<https://www.tensorflow.org/install/source#gpu>)

2. Install TensorFlow with GPU Support:

Use pip to install TensorFlow GPU version: `pip install tensorflow-gpu`


(<https://www.tensorflow.org/install/pip#macos>)

3. Verify GPU Access:

```
import tensorflow as tf
```

```
print("Num GPUs Available: ", len(tf.config.experimental.list_physical_devices('GPU')))
```

TensorFlow for Multi-GPU



1. Hardware Setup:

Ensure the machine has multiple NVIDIA GPUs installed (check `nvidia-smi`)

2. Software Setup:

Similar to the single GPU setup, ensure all drivers, CUDA, and cuDNN are installed, and then install TensorFlow GPU.

3. Enable Multi-GPU Support:

Use `tf.distribute.Strategy`, TensorFlow's standard way to distribute computations across multiple devices.

4. MirroredStrategy for Multi-GPU:

```
strategy = tf.distribute.MirroredStrategy()
```

5. Wrap Model Building and Compilation:


```
with strategy.scope():  
    model = .....# Build your model here  
    model.compile(...)
```

6. Batch Size Adjustment:

Increase batch size accordingly

7. Train the Model:

Train your model as usual. TensorFlow handles the distribution of computations.



Tips for Multi-GPU Training

- Data Loading: Make sure the data/ batch_size is efficient enough to supply to all GPUs
- Performance Tuning: Monitor the usage of GPUs (using nvidia-smi) to ensure balanced utilization
- Debugging: Start with a single GPU with shallow model
- while single GPU setups in TensorFlow are relatively straightforward

Here are the steps for running Jupyter Notebook and a Python script on a remote GPU server:



Generate SSH Key Pair: - ssh-keygen -b 2048 -t rsa

- Choose a path to save your SSH key pair and set a password when prompted.
- Install the private SSH key on your local computer

Connecting to the GPU Server and Running Jupyter Notebook:

1. SSH into the Remote Server (Local):

- ssh -i path/to/your/private/key user@remote_server_ip

2. Create a Screen Session (Remote):

- screen -S sessionname

3. Set GPU Visibility (Remote):

- export CUDA_VISIBLE_DEVICES=2,3

4. Start Jupyter Notebook (Remote):

- jupyter notebook --no-browser --port 1234

5. Detach from the Screen Session (Remote): - Press Ctrl + a followed by d to detach from the screen session

6. Reattach to the Screen Session (Local):


- screen -S sessionname

7. Set Up Port Forwarding (Local):

- ssh -NL 1234:localhost:1234 user@remote_server_ip

8. Detach from the Screen Session (Local):

- Press Ctrl + a followed by d to detach from the screen session



Thank you!!!

Questions?

sonal.rami@awi.de